Week 1 Day 2 – GenAI Home Assignment: 10 AI Model Providers

Model Provider	Model	Release Date	Knowledge Cutoff	Token Cost (Input / Output)	Context Window	Type	When to Use
OpenAI	GPT-5.4	5 Mar 2026	31 Aug 2025	$2.50 / $15 per 1M	1,050,000	Multimodal reasoning LLM	Complex reasoning, coding, document analysis and AI agents
Anthropic	Claude Opus 4.6	5 Feb 2026	May 2025	$5 / $25 per 1M	1M (beta at launch)	Multimodal reasoning LLM	Complex coding, debugging and large codebases
Google DeepMind	Gemini 3.1 Pro	19 Feb 2026	Jan 2025	$2 / $12 per 1M (up to 200K input)	1,048,576	Multimodal reasoning LLM	Video, audio, image and document analysis
Mistral AI	Mistral Small 4	16 Mar 2026	Not publicly confirmed	$0.15 / $0.60 per 1M	256K	Multimodal LLM	Cost-efficient chat, coding and document processing
DeepSeek	DeepSeek-V4-Pro	24 Apr 2026	Not publicly confirmed	$0.66 / $1.98 per 1M (off-peak)	1M	Reasoning and coding LLM	Complex coding, debugging and agent workflows
Amazon	Nova 2 Lite	2 Dec 2025	Oct 2025 (unconfirmed)	AWS region-dependent	1M	Multimodal reasoning LLM	Enterprise automation, support and documents
Cohere	Command A+	20 May 2026	1 Apr 2025	Free within API limits; deployment varies	128K	Multimodal MoE LLM	Enterprise agents and multilingual documents
xAI	Grok 4.3	2026 (exact date unverified)	Not publicly confirmed	$1.25 / $2.50 per 1M (short context)	1M	Multimodal reasoning LLM	Coding, reasoning and tool calling
Alibaba Cloud	Qwen3.5 Plus	20 Apr 2026 (snapshot)	Not publicly confirmed	Region- and tier-dependent	1M	Multimodal LLM	Coding agents, video and long-context analysis
IBM	Granite 4.1 8B	29 Apr 2026 (family)	Not publicly confirmed	Self-hosted; infrastructure costs	128K	Open-weight language model	Private enterprise AI, RAG and assistants

Key Parameters Explained

Input tokens: Units of text or other content sent to the model, such as your prompt, instructions and supplied documents.

Output tokens: Units generated by the model in its response, such as an answer, summary or code.

Cost of tokens: The price charged for processing input tokens and generating output tokens, typically shown per million tokens. Prices may vary by usage tier.

Context window: The maximum amount of information a model can hold in its active context, measured in tokens. It includes the prompt, conversation and generated output, subject to model-specific limits.

Knowledge cutoff: The approximate latest date covered by a model’s training knowledge, when the provider discloses it. This is different from the release date.

Model type: A description of what the model can do, such as text reasoning, coding or multimodal processing of images, audio and video.